

LECTURE 10

Yue Gong 2020.02.18

Review

Hash-based data structure

- Hash Tables / Maps / Sets
 - Exact equality lookup
- Hash-based similarity lookup
 - Minhash
- Approximate Membership Query (is q in a set)
 - Bloom Filter
- Approximate Counting Query (how many times has a q shown up)
 - Count-Min Sketch

- Approximate Distinct Count
 - Relations to Machine Learning
- } This Lecture

Distinct Count Problem

Estimate the number of distinct elements in a list with a single pass

Exact Solution:

$O(\# \text{ Distinct Elems})$ Memory


- What if we only have $O(m)$ space $m \ll D$?

hash function: $h: S \rightarrow \{0, \dots, 2^L - 1\}$

What is the probability a hash code is odd?

 $\frac{1}{2}$
$$h: S \rightarrow \{0, 1\}^L$$

What is the probability that the leftmost bit is 1?



x = the position of the least-significant 1

$$P(X=1) = \frac{1}{2}$$

$$P(X=2) = \frac{1}{4}$$

$$P(X=3) = \frac{1}{8}$$

$$P(x=L) = \frac{1}{2^L}, \text{ where } x \text{ is the position}$$

of the least-significant 1.

$$\text{BMP}[0, 1, \dots, L-1]$$

for s in S :

p = first index with a 1 from the left in the hash(s)

$$BMP[p] = 1$$

BMP $\sqcup \sqcup \sqcup \sqcup \dots \sqcup$
 $\frac{n}{2} \frac{n}{4} \dots \frac{n}{2^{i-1}}$

Estimation: Calculate $R \rightarrow$ right most bit that is 1

Then, estimate 2^R as the # distinct values

This is called "Flajolet-Martin Sketch"

Relations to Machine Learning

→ high-dimensional sparse features

one-hot embedding

A, B, C

A, B, C

1. A, B, C $\rightarrow [1, 1, 1] \rightarrow \begin{matrix} A/C & B \\ 1 & 1 \end{matrix}$

2. $A, C \rightarrow [1, 0, 1] \rightarrow [1 \ 0]$

3. $B \rightarrow [0, 1, 0] \rightarrow [0 \ 1]$

4. $A, C \rightarrow [1, 0, 1] \rightarrow [1, 0]$

5. $B \rightarrow [0, 1, 0] \rightarrow [0, 1]$

Are they compressible?

Yes! Because A always appears with C.

Feature Hashing

$$[v_1, v_2, v_3, \dots, v_j]$$
$$[v'_1, v'_2, \dots, v'_m]$$

Reduce Feature dimension

Feature Hashing (Yahoo Paper)

v = original feature vector $0, 1, \dots, i-1$

x = compressed feature vector $0, 1, \dots, m-1$

$$h_1 = \{0, \dots, i-1\} \rightarrow \{0, \dots, m-1\}$$

maps and randomly groups together features

$$h_2 = \{0, \dots, i-1\} \rightarrow \{-1, 1\}$$

Assign a negative one or positive one, to each one of the features

$$\rightarrow x[j] = \sum_{l=0}^{i-1} \delta(h_1(l)=j) \cdot v[l] \cdot \underbrace{h_2(l)}_{\substack{\downarrow \\ \text{randomly flips} \\ \text{the sign}}}$$

One Interesting Property

$$\langle v, v' \rangle = \sum_{l=0}^{i-1} v[l] \cdot v'[l]$$

$$\underline{E(\langle x, x' \rangle)} = \langle v, v' \rangle, \text{ feature hashing preserves the inner product between examples.}$$

why inner product can be preserved?

$$\langle x, x' \rangle = \sum_{l=0}^{i-1} \sum_{m=0}^{i-1} \delta(h_1(l)=h_1(m)) \cdot \underbrace{h_2(l)h_2(m)}_{\substack{\downarrow \\ \text{add together}}} \cdot v[l]v[m]$$

add together

$$h_2(l), h_2(m) \text{ have the same sign} \rightarrow h_2(l)h_2(m) = 1$$

$$h_2(l), h_2(m) \text{ have different signs} \rightarrow h_2(l)h_2(m) = -1$$

↓ cancellation

Why is preserving inner product so good?

$$\min_{\theta} \|x^T \theta - y\|_2^2$$

$$\text{if } x \text{ is full rank, } \theta = \underbrace{(x^T x)^{-1}}_{\substack{\uparrow \\ \text{Matrix of Inner Product}}} x^T y$$

Matrix of Inner Product

Advantages over PCA

If using PCA, needs to materialize a dictionary of words

→ if you use Feature Hashing, you can skip materialization

e.g. { the quick brown }

hash ↓ hash ↓ hash ↓